
Introduction to Recommendation Engine

— Ching Lee —
2018/09/07





Prologue

- People are increasing the reliance on conveniences such as e-commerce store or streaming entertainment.
- “Guess” what the customers may like in advance, so to promote more things to sell and in turn generate more revenue.
- A recommendation engine (RE) is any kind of model that can infer the relationship between users/items and make proper prediction for the users.









Gender	Male
Age	30
Prefer genre	Sci-fi, comedy, action
Prefer director	Steven Spielberg, Christopher Nolan, Michael Bay, James Cameron
Prefer actor/actress	Tom Hanks, Leonardo DiCaprio, Anne Hathaway, Scarlet Johansson





Gender	Male
Age	30
Prefer genre	Sci-fi, comedy, action
Prefer director	Steven Spielberg, Christopher Nolan, Michael Bay, James Cameron
Prefer actor/actress	Tom Hanks, Leonardo DiCaprio, Anne Hathaway, Scarlet Johansson











Content Filtering

- What we have just conducted is essentially one way of doing recommendation: *content filtering*.
- By **building profiles** for both the users and movies, we can provide recommendation by matching the **content** between the two groups, hence the name.
- Take a lot of efforts to build such profiles and many conditional settings to fit a person's taste. 🤪

Strategies for Recommendation

- Beside *content filtering* we had just mentioned, there is another method called *collaborative filtering (CF)*.
- *CF* relies on past user behavior among a group of users (hence *collaborative*), so we aren't required to create profiles explicitly.
- The two primary areas of *CF* are the *neighborhood methods (memory-based)* and *latent factor models*.





User-based and Item-based CF


- For neighborhood methods, we have two primary types of algorithms: **user-based** and **item-based**.
- In *user-based CF*, we group together users who gave similar ratings to the same set of items, whereby we could later use the ratings of a specific users to predict those of his/her peers.
- For *item-based CF*, we group the items instead.



8.7	10	10	9	?	?	?	9	8	8.5
?	?	?	1	10	10	10	?	?	2
8	9	9	9	?	?	5	10	10	9.5

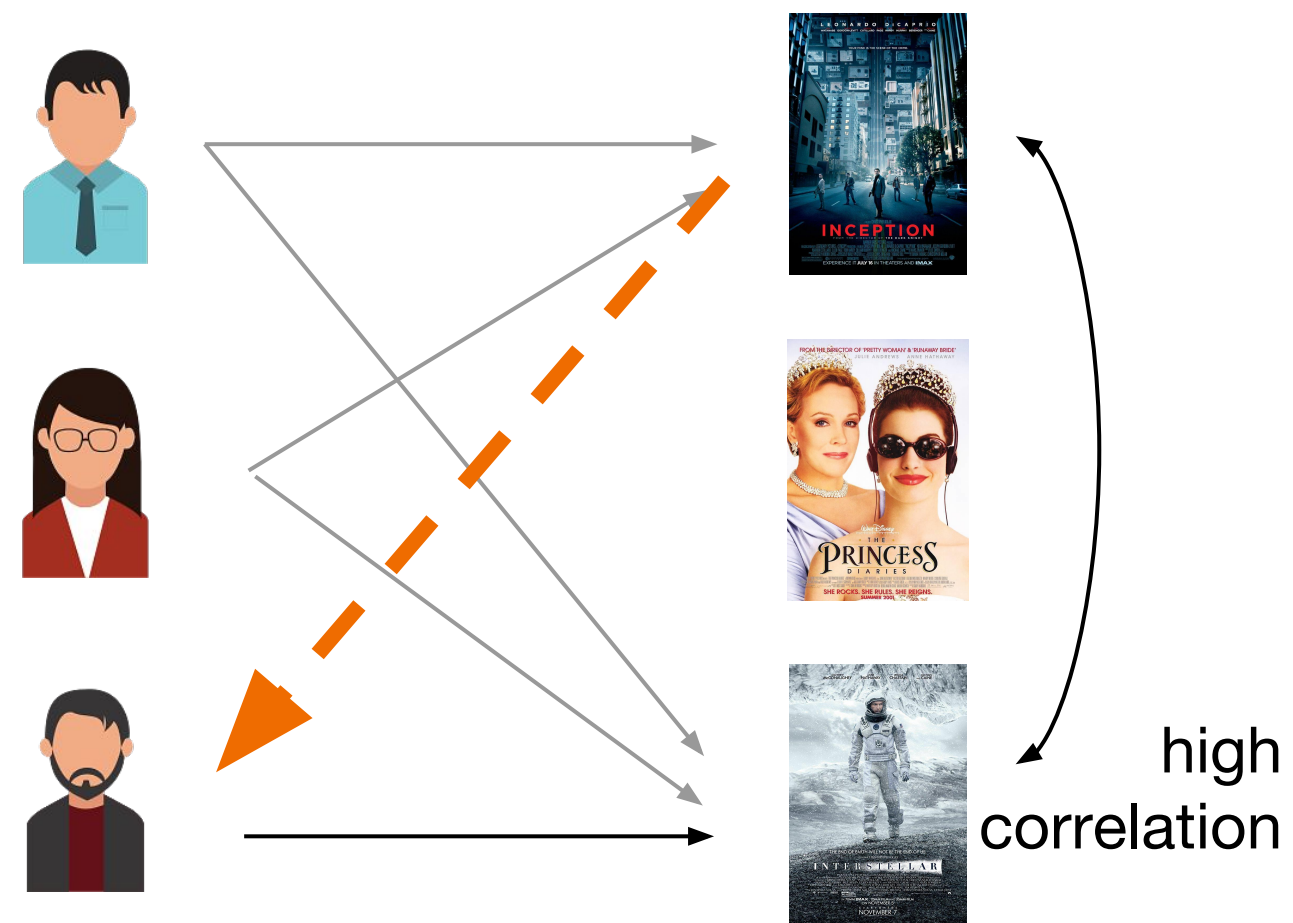
--- User-based CF

--- Item-based CF



User-based filtering

You may like it because your “friends” liked it.




Item-based filtering

You tend to like that item since you have liked those items.

Measuring Similarity

- To measure the similarity between users or items, we can use metrics like **Euclidean distance** or **cosine similarity**.










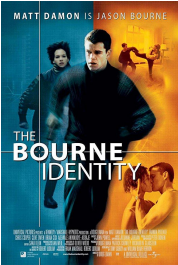
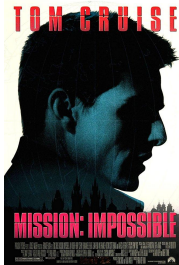



Euclidean distance: distance between points

$$d(\mathbf{p}, \mathbf{q}) = d(\mathbf{q}, \mathbf{p}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \cdots + (q_n - p_n)^2}$$

Cosine similarity: angle between vectors

$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\|_2 \|\mathbf{B}\|_2} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

Example: User-Based CF

											
	8.7	10	10	9	?	?	?	9	8	8.5	23.96
	?	?	?	1	10	10	10	?	?	2	17.46
	8	9	9	9	?	?	5	10	10	9.5	24.94


→ cosine similarity:

$$\text{User 1} \text{ and } \text{User 2} : \frac{9 * 1 + 8.5 * 2}{23.96 * 17.46} = \mathbf{0.062}$$

$$\text{User 1} \text{ and } \text{User 3} : \frac{8.7 * 8 + 10 * 9 + 10 * 9 + 9 * 9 + 9 * 10 + 8 * 10 + 8.5 * 9.5}{23.96 * 24.94} = \mathbf{0.973}$$




Latent Factor Models

- Explain the rating by characterizing both users and items on a number of *latent factors* inferred from rating patterns.



The factors are latent.

No one
cares ...

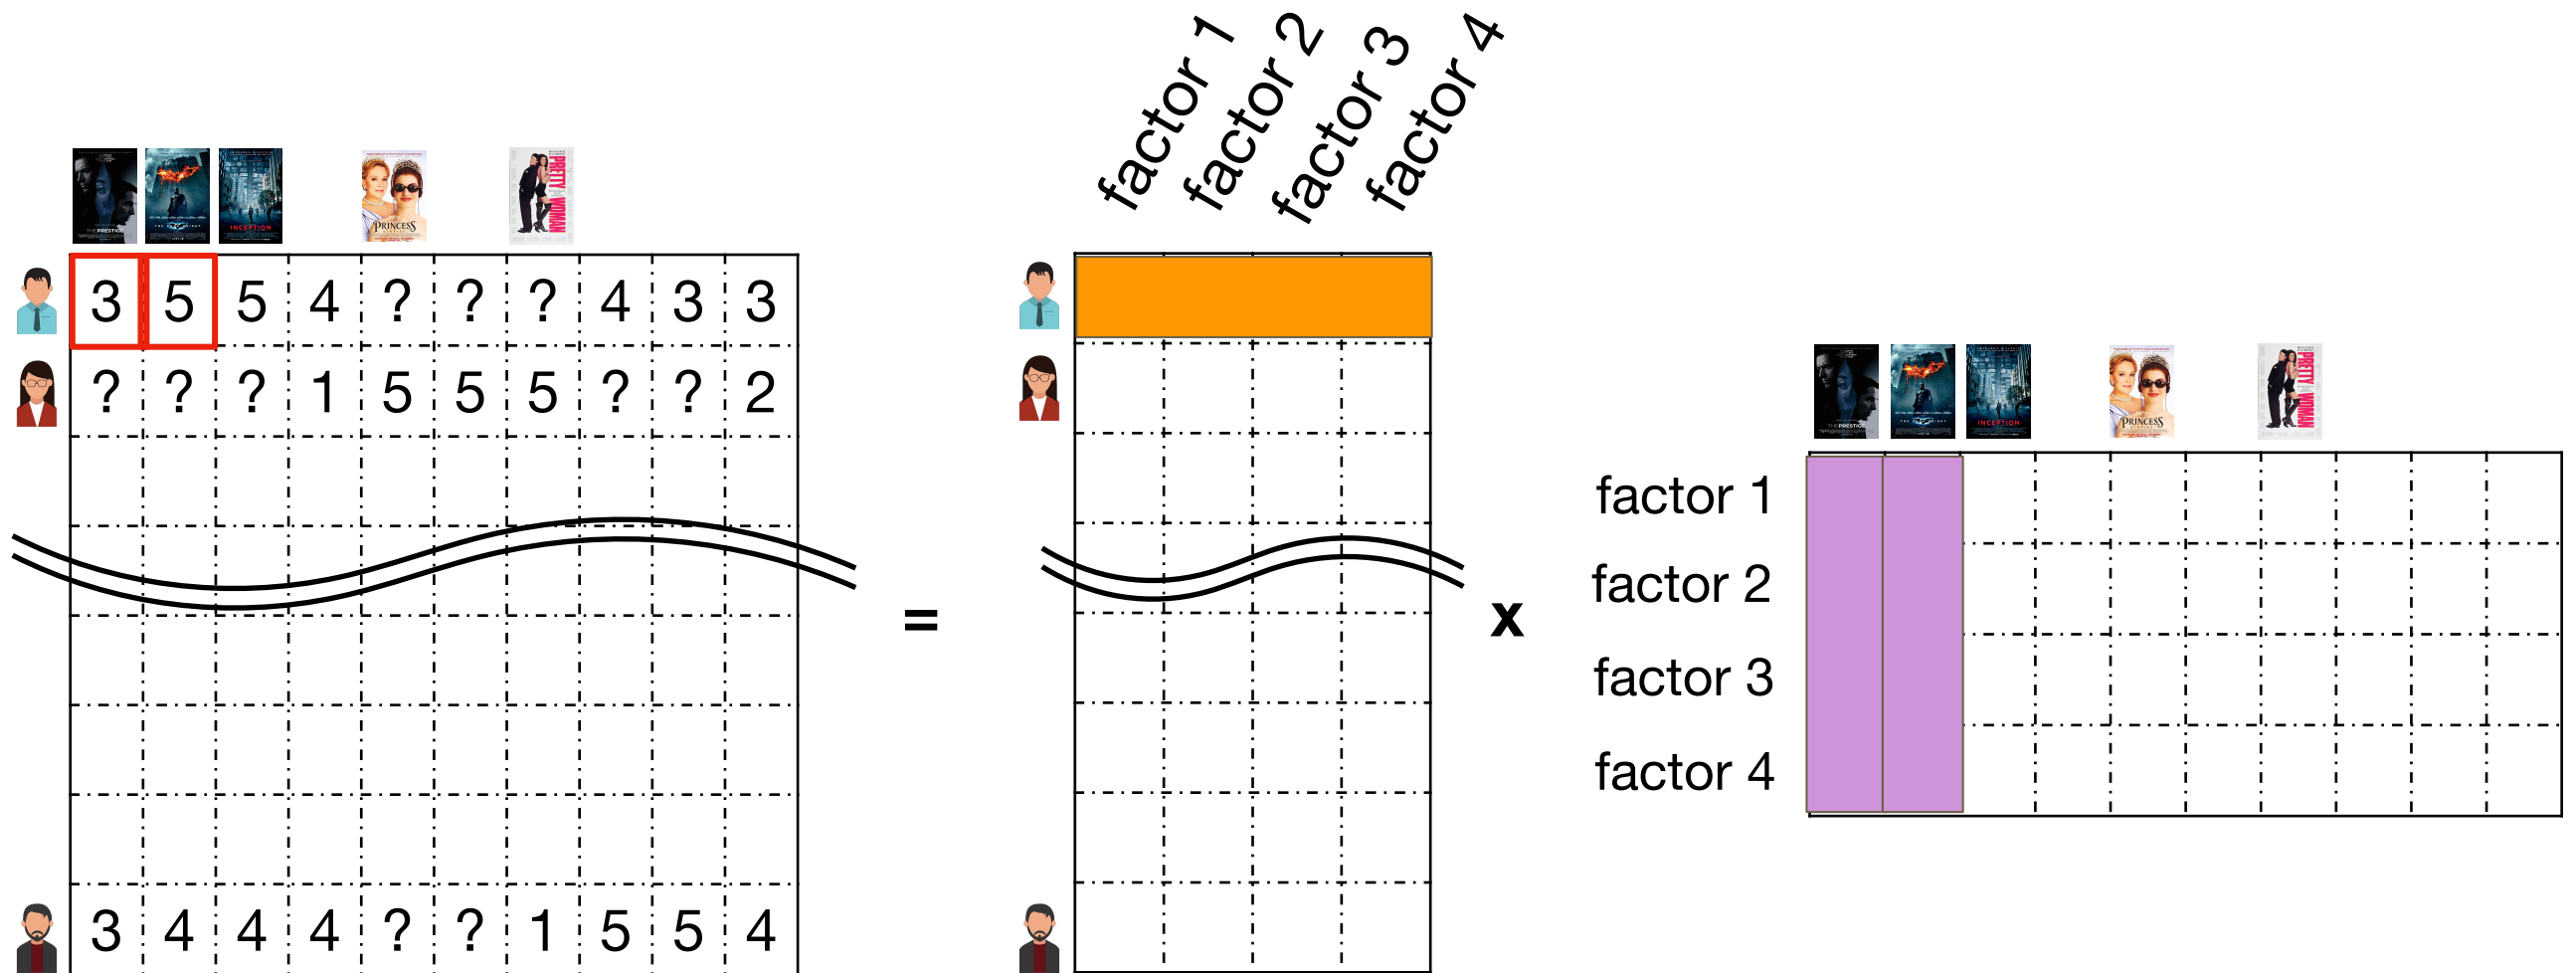


Not directly
observable ...

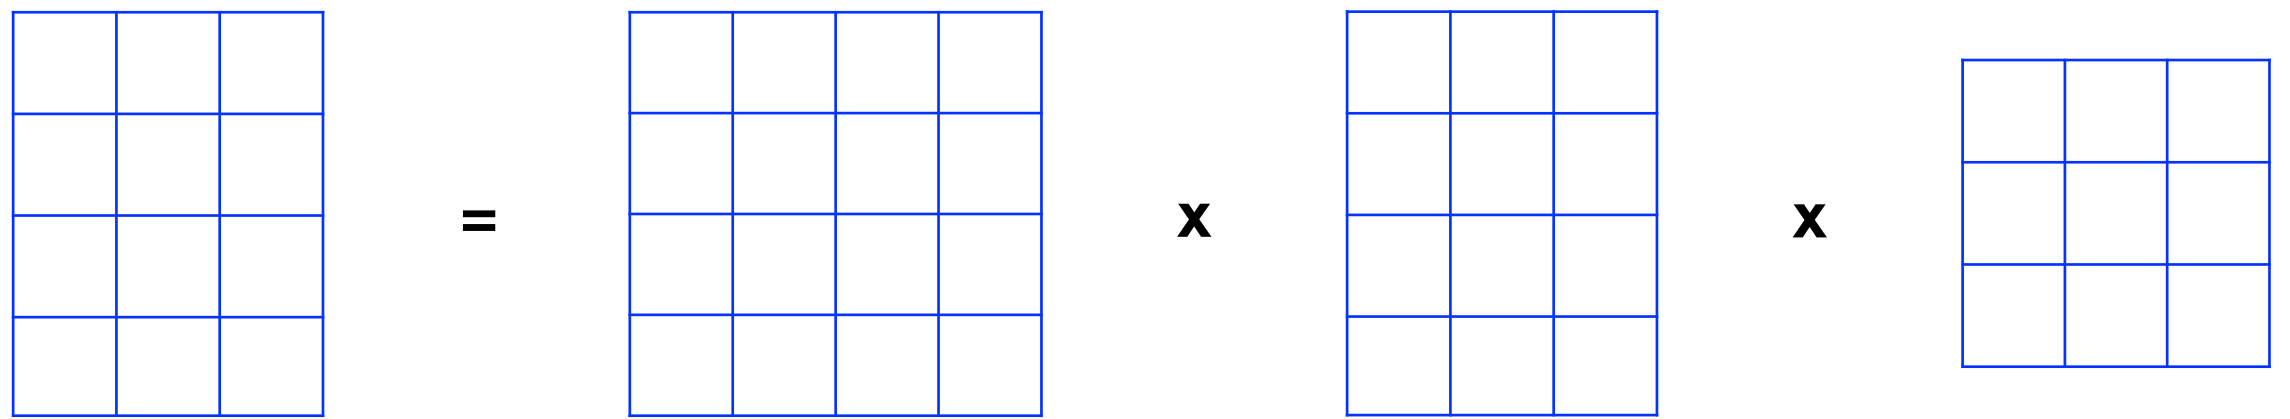


3	5	5	4	?	?	?	4	3	3
?	?	?	1	5	5	5	?	?	2
3	4	4	4	?	?	1	5	5	4

What is latent factor anyway?




Singular Value Decomposition (SVD)




The diagram illustrates the Singular Value Decomposition (SVD) of a matrix A . It shows a 4x3 grid representing matrix A on the left, followed by an equals sign. To the right of the equals sign is a 4x4 grid representing matrix U , followed by a multiplication symbol 'x'. This is followed by a 4x3 grid representing matrix Σ , followed by another multiplication symbol 'x'. Finally, on the right, is a 3x3 grid representing matrix V^T . Below each grid is its corresponding mathematical symbol: A , U , Σ , and V^T .

$$A = U \Sigma V^T$$

That is (kind of) what SVD is



But in practice...



3	5	5	4	?	?	?	4	3	3
?	?	?	1	5	5	5	?	?	2
3	4	4	4	?	?	1	5	5	4


“Full of holes”

=

	factor 1									
	factor 2									
	factor 3									
	factor 4									

x

factor 1
factor 2
factor 3
factor 4






Can't change the length of “contact”


How to factor the matrix

- To perform matrix factorization for large matrices, we learn the entries through optimization methods such as *stochastic gradient descent* (SGD).
- Methods like *alternating least square* (ALS) are also used when computation can be parallelized.
- We are going to briefly introduce SGD for its popularity.


Goal:



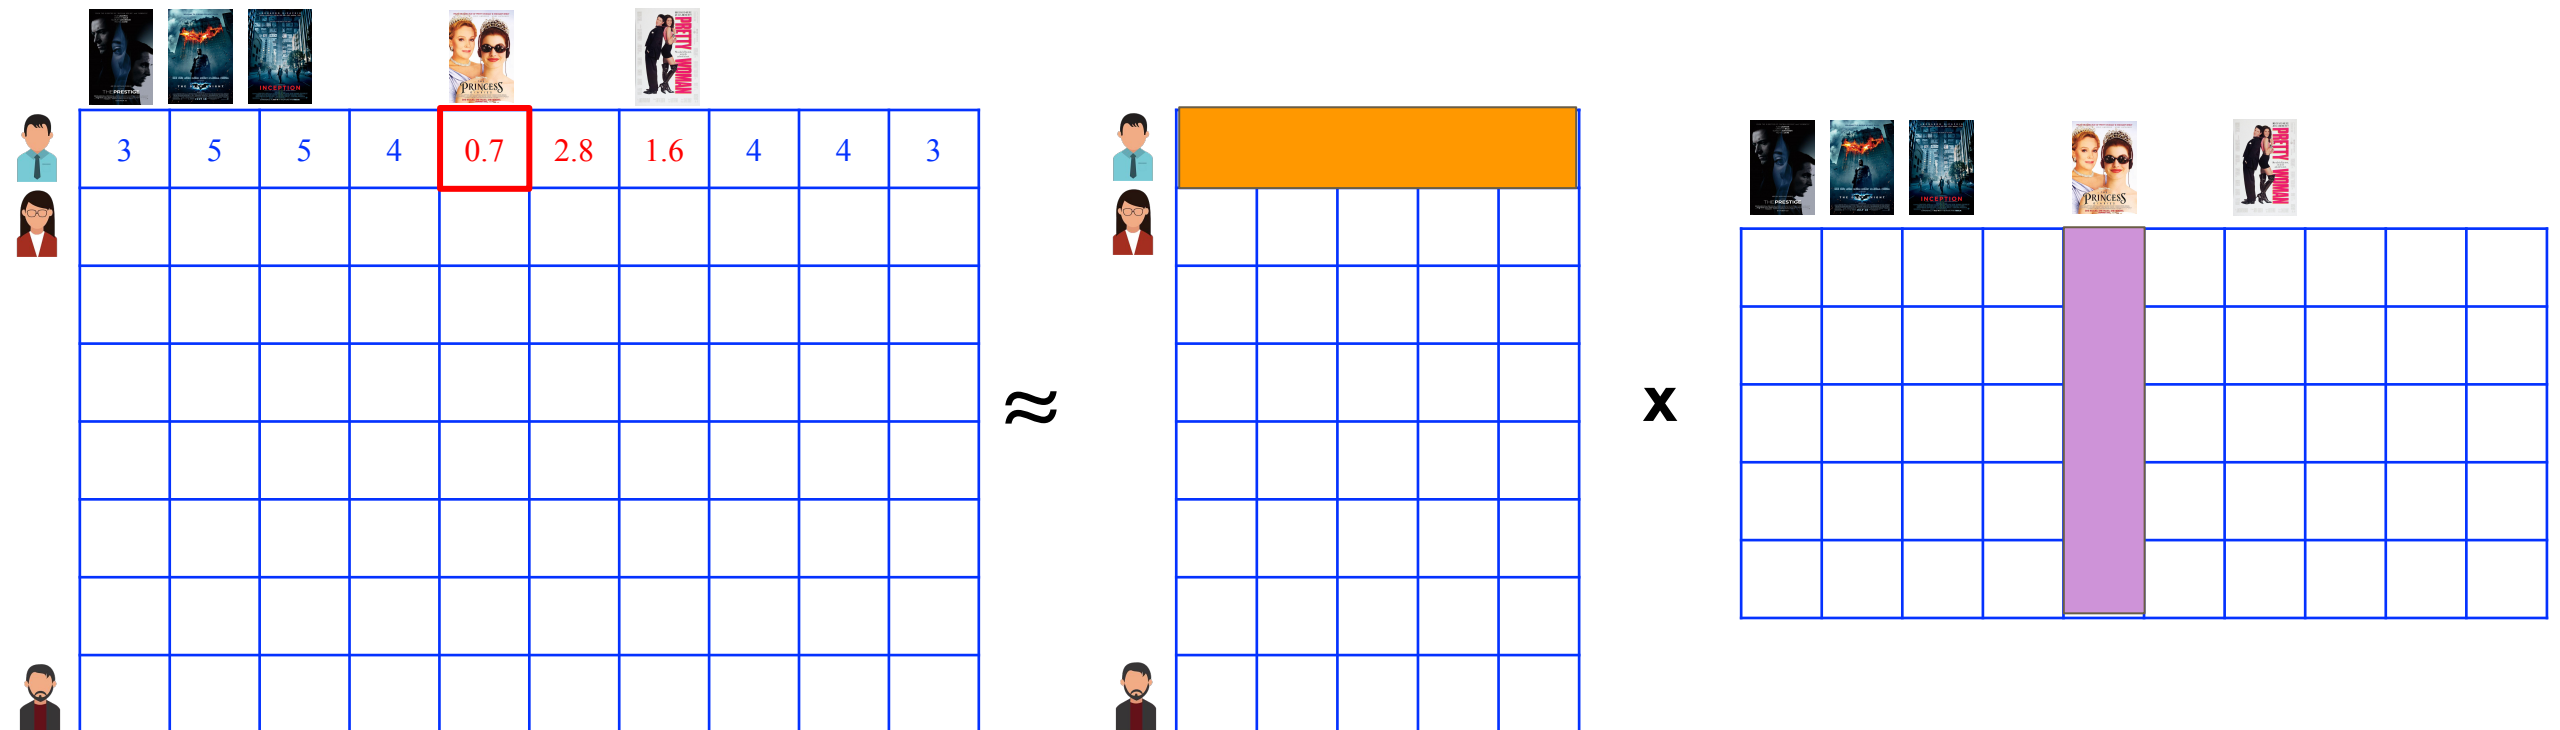
Initialize these and tune the numbers!



x



Ok, I've got the matrices. Then what?



→ Key points:

1. The approximate matrix will **not** be identical to the original.
2. The factor matrices will keep changing as long as there are users changing the rating (even if *you* stay inactive for a while).

Evaluating a model

1. Root Mean Square Error / Mean Absolute Error

$$\sqrt{\sum_i (Pred_i - True_i)^2} \quad \sum_i |Pred_i - True_i|$$

2. Confusion Matrix (Precision and Recall)

		Actual	
		Positive	Negative
Predicted	Positive	True Positive	False Positive
	Negative	False Negative	True Negative

3. Discounted Cumulative Gain (DCG)

Evaluating a model

- To evaluate how a machine learning model did, we use metrics such as *precision* and *recall*.

		Actual value	
		True	False
Predicted value	True	True positive (TP)	False positive (FP)
	False	False negative (FN)	True negative (TN)

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Scenario 1:

		Actual value	
		Have cancer	Safe
Predicted value	Have cancer	45	5
	Safe	190	1000

$$\text{Precision} = \frac{45}{45 + 5} = 90\%$$

$$\text{Recall} = \frac{45}{45 + 190} = 19.15\%$$

Scenario 2:

		Actual value	
		Have cancer	Safe
Predicted value	Have cancer	200	800
	Safe	35	205

$$\text{Precision} = \frac{200}{200 + 800} = 20\%$$

$$\text{Recall} = \frac{200}{200 + 35} = 85.11\%$$

That's why we have F1-score

- F1-score (or F-score) is the harmonic mean between precision and recall:


$$\text{F1} = 2 * \frac{1}{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}} = 2 * \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

Receiver Operating Characteristic (ROC) curve


		Actual value	
		True	False
Predicted value	True	True positive (TP)	False positive (FP)
	False	False negative (FN)	True negative (TN)

True positive rate: $TPR = TP / (TP + FN)$

False positive rate: $FPR = FP / (FP + TN)$



Let's evaluate the ranking as well

- A RE often delivers numerous outputs, while only a portion of them is most relevant to what the user really wants.
- We **rank our results** for the users, so the entries that the user would most likely selected would be near the top.
- How do we evaluate the *ranking* of results? 



(鍾欣桐)

1



2



3



(鍾欣凌)

4

Relevance (0-3 scale):



1

Cumulative Gain: $1 + 3 + 3 + 1 = 8$
(@ rank 4)

Discounted Cumulative Gain (DCG):
(@ rank p)

$$\sum_{i=1}^p \frac{rel_i}{\log_2(i+1)}$$


3

$$\frac{1}{\log_2(1+1)} + \frac{3}{\log_2(2+1)} + \frac{3}{\log_2(3+1)} + \frac{1}{\log_2(4+1)} = 4.824$$



3

$$\frac{3}{\log_2(1+1)} + \frac{3}{\log_2(2+1)} + \frac{1}{\log_2(3+1)} + \frac{1}{\log_2(4+1)} = 5.824$$



IDCG : 最理想化的分數



1

Normalized DCG:
(@ rank p)

$$\frac{DCG_p}{IDCG_p}$$